

```
Question 1

In [2]: import cv2
import numpy as np

In [3]: import math

In [4]: import matplotlib.pyplot as plt

In [5]: xray_img = cv2.imread("xray.png",-1)
xray_img.shape
Out[5]: (1024, 1024)

In [34]: plt.imshow(xray_img)
plt.title("Original Image")
plt.axis('off')
Out[34]: (-0.5, 1023.5, 1023.5, -0.5)

Original Image


In [56]: noisy_img = cv2.imread("xray_sp.png",-1)
plt.imshow(noisy_img,cmap='gray')
plt.title("Noisy Image")
plt.axis('off')
Out[56]: (-0.5, 1023.5, 1023.5, -0.5)

Noisy Image


In [182]: median_filtered_img = cv2.medianBlur(noisy_img,3)
plt.imshow(median_filtered_img,cmap='gray')
plt.axis('off')
Out[182]: (-0.5, 1023.5, 1023.5, -0.5)



In [37]: cv2.imwrite("median_filtered_img.png",median_filtered_img)
Out[37]: True

In [48]: def rmse(img1,img2):
    mse = np.square(np.subtract(img1,img2)).mean()
    rmse = math.sqrt(mse)
    return rmse

In [46]: rmse_median_filter = rmse(xray_img,median_filtered_img)

In [50]: print(f"RMSE of Median Filtered Image: {rmse_median_filter}")
RMSE of Median Filtered Image: 1.017037216708072

In [49]: def snr(filtered_img,original_img):
    num = np.square(filtered_img).sum()
    diff_sq = np.square(filtered_img - original_img)
    denom = np.sum(diff_sq)
    snr = num/denom
    return snr

In [160]: snr_median_filter = snr(median_filtered_img,xray_img)

In [161]: print(f"SNR of Median Filtered Image: {snr_median_filter}")
SNR of Median Filtered Image: 104.28556078045742

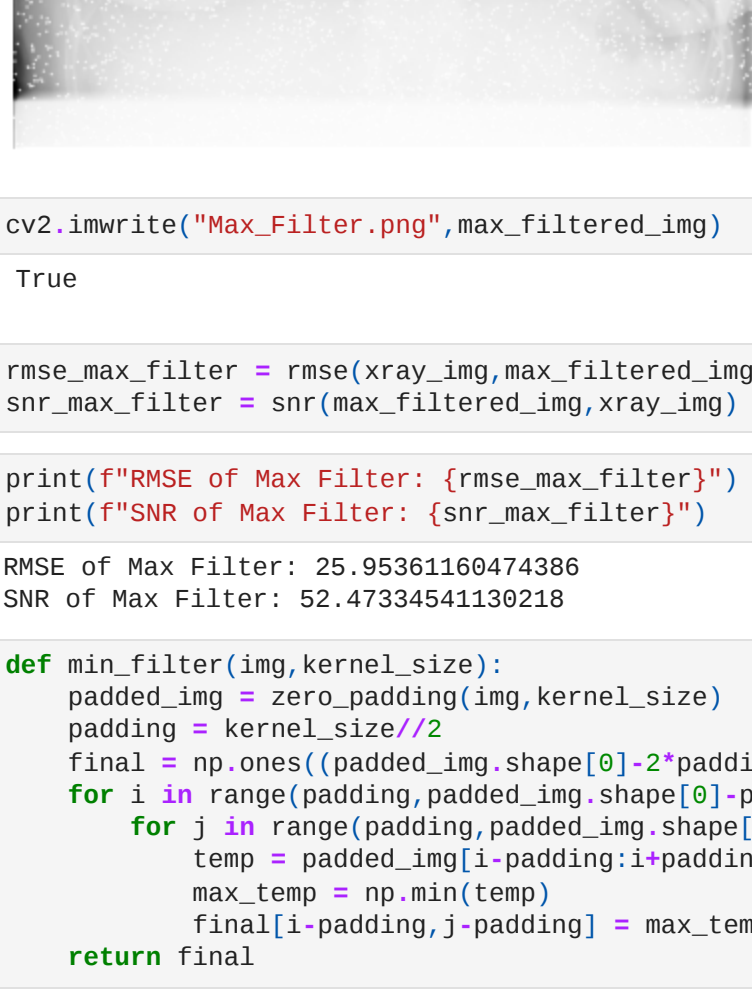
In [98]: def zero_padding(img,kernel_size):
    h,w = img.shape
    padding = kernel_size/2
    new_h = h+padding*2
    new_w = w+padding*2
    temp = np.zeros((new_h,new_w))
    for i in range(h):
        for j in range(w):
            temp[i+padding,j+padding] = img[i,j]
    return temp

In [198]: def max_filter(img,kernel_size):
    padded_img = zero_padding(img,kernel_size)
    padding = kernel_size/2
    final = np.ones((padded_img.shape[0]-2*padding,padded_img.shape[1]-2*padding))
    for i in range(padding,padded_img.shape[0]-padding):
        for j in range(padding,padded_img.shape[1]-padding):
            #i,j is the centre pixel
            temp = padded_img[i-padding:i+padding+1,j-padding:j+padding+1]
            max_temp = np.max(temp)
            final[i-padding,j-padding] = max_temp
    return final

In [199]: max_filtered_img = max_filter(noisy_img,3)

In [200]: max_filtered_img.shape
Out[200]: (1024, 1024)

In [261]: plt.imshow(max_filtered_img,cmap="gray")
plt.axis('off')
Out[261]: (-0.5, 1023.5, 1023.5, -0.5)



In [202]: cv2.imwrite("Max_Filter.png",max_filtered_img)
Out[202]: True


In [203]: rmse_max_filter = rmse(xray_img,max_filtered_img)
snr_max_filter = snr(max_filtered_img,xray_img)

In [204]: print(f"RMSE of Max Filter: {rmse_max_filter}")
print(f"SNR of Max Filter: {snr_max_filter}")
RMSE of Max Filter: 25.95261160474306
SNR of Max Filter: 52.47334541139218

In [205]: def min_filter(img,kernel_size):
    padded_img = zero_padding(img,kernel_size)
    padding = kernel_size/2
    final = np.ones((padded_img.shape[0]-2*padding,padded_img.shape[1]-2*padding))
    for i in range(padding,padded_img.shape[0]-padding):
        for j in range(padding,padded_img.shape[1]-padding):
            #i,j is the centre pixel
            temp = padded_img[i-padding:i+padding+1,j-padding:j+padding+1]
            max_temp = np.min(temp)
            final[i-padding,j-padding] = max_temp
    return final

In [206]: min_filtered_img = min_filter(noisy_img, 3)

In [207]: plt.imshow(min_filtered_img,cmap="gray")
plt.axis('off')
Out[207]: (-0.5, 1023.5, 1023.5, -0.5)

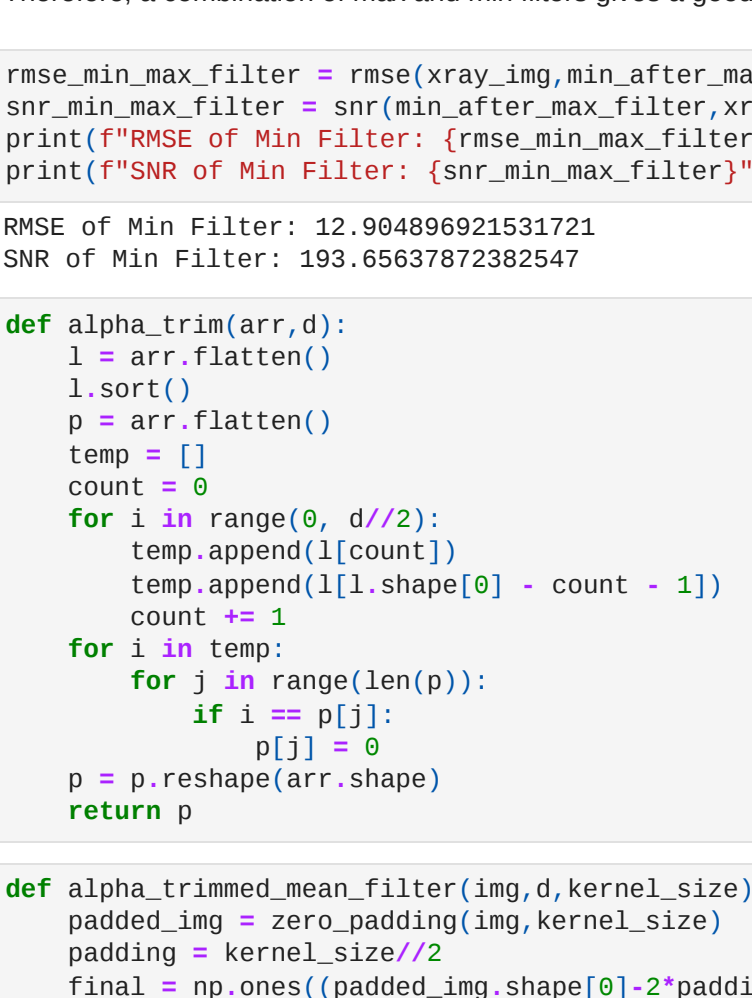


In [208]: cv2.imwrite("Min_Filter.png",min_filtered_img)
Out[208]: True

In [209]: rmse_min_filter = rmse(xray_img,min_filtered_img)
snr_min_filter = snr(min_filtered_img,xray_img)
print(f"RMSE of Min Filter: {rmse_min_filter}")
print(f"SNR of Min Filter: {snr_min_filter}")
RMSE of Min Filter: 20.75023186795397
SNR of Min Filter: 73.26402100026705

In [218]: min_after_max_filter = min_filter(max_filtered_img,5)

In [219]: plt.imshow(min_after_max_filter,cmap="gray")
plt.axis('off')
Out[219]: (-0.5, 1023.5, 1023.5, -0.5)



In [220]: cv2.imwrite("Min Filter on Max Filtered Image.png",min_after_max_filter)
Out[220]: True

Therefore, a combination of max and min filters gives a good result for salt and pepper noise

In [222]: rmse_min_max_filter = rmse(xray_img,min_after_max_filter)
snr_min_max_filter = snr(min_after_max_filter,xray_img)
print(f"RMSE of Min Filter: {rmse_min_max_filter}")
print(f"SNR of Min Filter: {snr_min_max_filter}")
RMSE of Min Filter: 32.904896021531721
SNR of Min Filter: 193.65037872302547

In [54]: def alpha_trim(arr,d):
    l = arr.flatten()
    l = sorted(l)
    p = arr.flatten()
    count = 0
    for i in range(0, d/2):
        temp.append(l[count])
        temp.append(l[l.shape[0] - count - 1])
        count += 1
    for i in temp:
        for j in range(len(p)):
            if i == p[j]:
                p[j] = 0
    p = p.reshape(arr.shape)
    return p

In [80]: alpha_trimmed_mean_filter(img,d,kernel_size):
    padded_img = zero_padding(img,kernel_size)
    padding = kernel_size/2
    final = np.ones((padded_img.shape[0]-2*padding,padded_img.shape[1]-2*padding))
    for i in range(padding,padded_img.shape[0]-padding):
        for j in range(padding,padded_img.shape[1]-padding):
            #i,j is the centre pixel
            temp = padded_img[i-padding:i+padding+1,j-padding:j+padding+1]
            alpha_trim_temp = alpha_trim(temp,d)
            final[i-padding,j-padding] = alpha_trim_temp.sum()/(kernel_size**2 - d)
    return final

In [80]: alpha_trimmed_xray_img = alpha_trimmed_mean_filter(noisy_img,0,3)

In [81]: alpha_trimmed_xray_img = cv2.normalize(alpha_trimmed_xray_img,None,0,255,cv2.NORM_MINMAX,cv2.CV_BUC1)

In [62]: plt.imshow(alpha_trimmed_xray_img,cmap='gray')
plt.axis('off')
Out[62]: (-0.5, 1023.5, 1023.5, -0.5)



In [63]: cv2.imwrite("Alpha Trimmed Mean Filter.png",alpha_trimmed_xray_img)
Out[63]: True

In [64]: rmse_alpha_filter = rmse(xray_img,alpha_trimmed_xray_img)
snr_alpha_filter = snr(alpha_trimmed_xray_img,xray_img)
print(f"RMSE of Min Filter: {rmse_alpha_filter}")
print(f"SNR of Min Filter: {snr_alpha_filter}")
RMSE of Min Filter: 2.8910151009102607
SNR of Min Filter: 11.87320020790404

Question 2

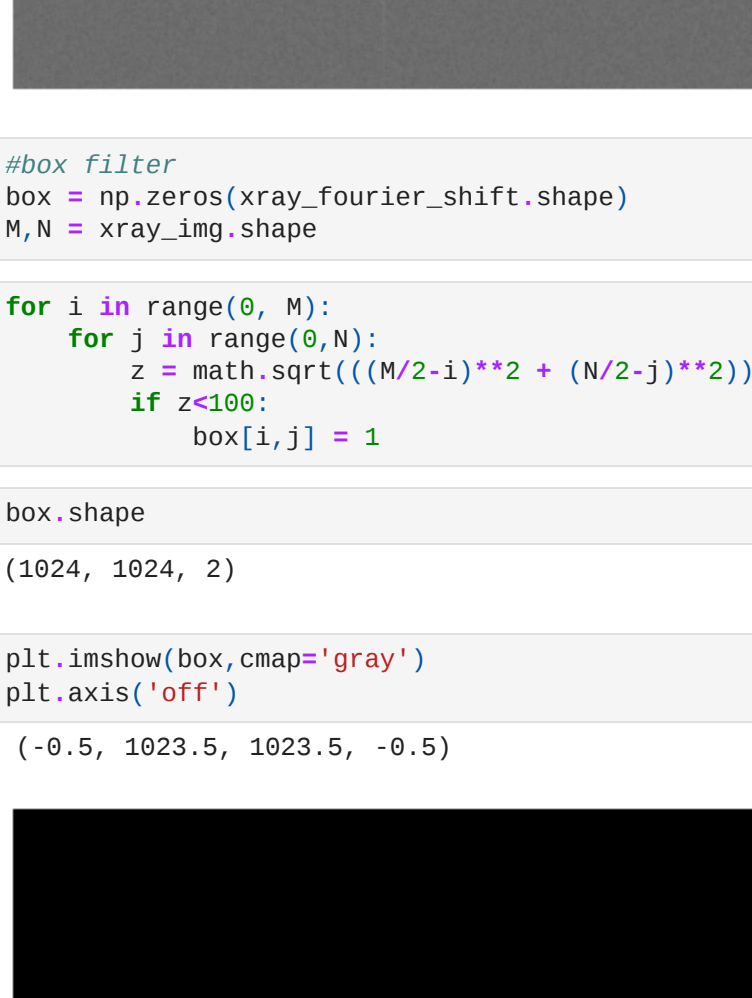
In [9]: noisy = cv2.imread("xray_g.png",-1)

In [10]: noisy_float = np.float64(noisy)
xray_fourier = cv2.dft(noisy_float,flags = cv2.DFT_COMPLEX_OUTPUT)
xray_fourier_shift.shape
Out[11]: (1024, 1024, 2)

In [12]: magnitude = 20*np.log(cv2.magnitude(xray_fourier_shift[:,0],xray_fourier_shift[:,1]))

In [13]: magnitude = cv2.normalize(magnitude,None,0,255,cv2.NORM_MINMAX,cv2.CV_BUC1)

In [14]: plt.imshow(magnitude,cmap='gray')
plt.axis('off')
Out[14]: (-0.5, 1023.5, 1023.5, -0.5)



In [16]: #box filter
box = np.zeros(xray_fourier_shift.shape)
M,N = xray_img.shape

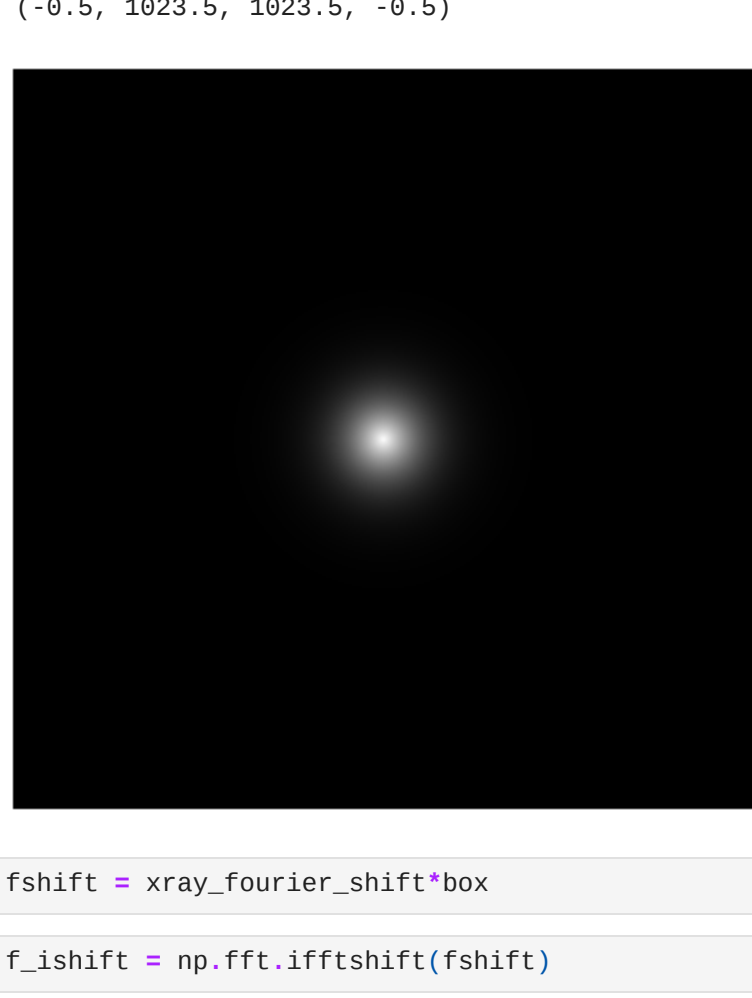
In [16]: for i in range(0, M):
    for j in range(0, N):
        z = math.sqrt(((N/2-i)**2 + (N/2-j)**2))
        if z<100:
            box[i,j] = 1

In [17]: box.shape
Out[17]: (1024, 1024, 2)

In [215]: plt.imshow(box,cmap='gray')
plt.axis('off')
Out[315]: (-0.5, 1023.5, 1023.5, -0.5)



In [358]: plt.imshow(gaussian,cmap='gray')
plt.axis('off')
Out[358]: (-0.5, 1023.5, 1023.5, -0.5)



In [43]: #Butterworth LPF
d0 = 300
n = 3
butter = np.zeros(xray_fourier_shift.shape)
for i in range(0, M):
    for j in range(0, N):
        z = math.sqrt(((N/2-i)**2 + (N/2-j)**2))
        butter[i,j] = 1/(1+math.exp(z/d0)**(2*n))

In [364]: plt.imshow(butter,cmap='gray')
plt.axis('off')
Out[364]: (-0.5, 1023.5, 1023.5, -0.5)



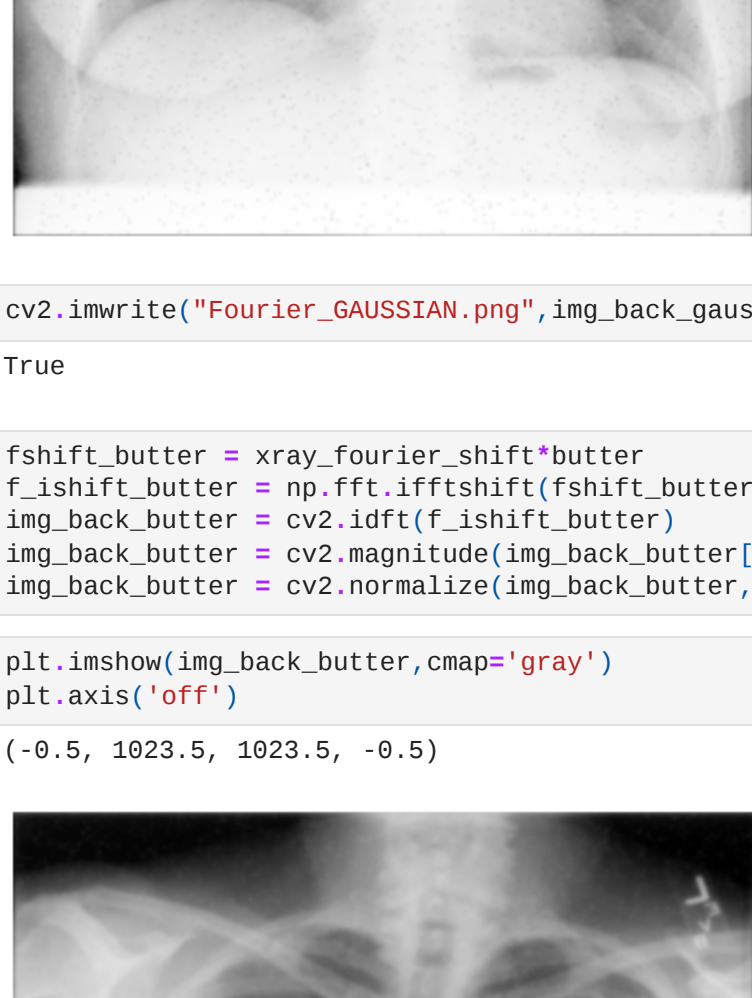
In [20]: fshift = xray_fourier_shift*box

In [21]: f_shift = np.fft.ifftshift(fshift)

In [23]: img_back = cv2.magnitude(img_back[:,0],img_back[:,1])

In [29]: img_back = cv2.normalize(img_back,None,0,255,cv2.NORM_MINMAX,cv2.CV_BUC1)

In [30]: plt.imshow(img_back,cmap='gray')
plt.axis('off')
Out[30]: (-0.5, 1023.5, 1023.5, -0.5)



In [31]: cv2.imwrite("Fourier_BOX.png",img_back)
Out[31]: True

In [34]: fshift_gaussian = xray_fourier_shift*gaussian
f_shift_gaussian = np.fft.ifftshift(fshift_gaussian)
img_back_gaussian = cv2.magnitude(img_back_gaussian[:,0],img_back_gaussian[:,1])
img_back_gaussian = cv2.normalize(img_back_gaussian,None,0,255,cv2.NORM_MINMAX,cv2.CV_BUC1)

In [35]: plt.imshow(img_back_gaussian,cmap="gray")
plt.axis('off')
Out[35]: (-0.5, 1023.5, 1023.5, -0.5)



In [36]: cv2.imwrite("Fourier_GAUSSIAN.png",img_back_gaussian)
Out[36]: True

In [44]: f_shift_butter = xray_fourier_shift*butter
f_shift_butter = np.fft.ifftshift(fshift_butter)
img_back_butter = cv2.magnitude(f_shift_butter)
img_back_butter = cv2.normalize(img_back_butter[:,0],img_back_butter[:,1])
img_back_butter = cv2.normalize(img_back_butter,None,0,255,cv2.NORM_MINMAX,cv2.CV_BUC1)

In [45]: plt.imshow(img_back_butter,cmap="gray")
plt.axis('off')
Out[45]: (-0.5, 1023.5, 1023.5, -0.5)



In [46]: cv2.imwrite("Fourier_BUTTERWORTH.png",img_back_butter)
Out[46]: True

RMSE and SNR values
rmse_box = rmse(xray_img,img_back)
rmse_gaussian = rmse(xray_img,img_back_gaussian)
rmse_butter = rmse(xray_img,img_back_butter)

In [52]: snr_box = snr(img_back,xray_img)
snr_gaussian = snr(img_back_gaussian,xray_img)
snr_butter = snr(img_back_butter,xray_img)

In [53]: print(f"Box Filter: RMSE = {rmse_box} | SNR = {snr_box}")
print(f"Gaussian Filter: RMSE = {rmse_gaussian} | SNR = {snr_gaussian}")
print(f"Butterworth Filter: RMSE = {rmse_butter} | SNR = {snr_butter}")
Box Filter: RMSE = 10.24232544042047 | SNR = 1.033821250909412
Gaussian Filter: RMSE = 4.0221692024322 | SNR = 0.8052044102752295
Butterworth Filter: RMSE = 4.275206706978054 | SNR = 5.871659126538234

In [ ]:
```



```
In [1]: import cv2
import numpy as np

In [2]: import math

In [3]: import matplotlib.pyplot as plt

In [4]: xray = cv2.imread('xray.png',-1)

In [5]: def zigzag(matrix):
    # Get the number of rows and columns in the matrix
    rows = len(matrix)
    cols = len(matrix[0])

    # Create an empty list to store the zigzag ordering
    zigzag = []
    temp = np.zeros((1,rows*cols))

    # Initialize the starting row and column
    row, col = 0, 0

    # Iterate over each diagonal of the matrix
    for i in range(rows + cols - 1):
        # If the diagonal is even, traverse it from top-right to bottom-left
        if i % 2 == 0:
            while row >= 0 and col < cols:
                zigzag.append(matrix[row][col])
                row -= 1
                col += 1
            if row < 0 and col <= cols - 1:
                row = 0
            if col <= cols:
                row += 2
                col -= 1
            # If the diagonal is odd, traverse it from bottom-left to top-right
        else:
            while col >= 0 and row < rows:
                zigzag.append(matrix[row][col])
                row += 1
                col -= 1
            if col < 0 and row <= rows - 1:
                col = 0
            if row <= rows:
                col += 2
                row -= 1
    zigzag = np.array(zigzag)
    zigzag = np.reshape(zigzag, (-1,matrix.shape[0]))

    return zigzag

In [6]: def reverse_zigzag_for8x8(matrix):
    for i in range(135):
        matrix= zigzag(matrix)
    return matrix

In [7]: def run_length_encode(input_string):
    # Initialize the encoded string and the counter
    encoded_string = ''
    count = 1

    # Iterate over each character in the input string
    for i in range(1, len(input_string)):
        if input_string[i] == input_string[i-1]:
            # If the current character is the same as the previous one, increment the counter
            count += 1
        else:
            # If the current character is different from the previous one, add the count and the character to the encoded string
            count_str = str(count)
            encoded_string += count_str + "," + str(input_string[i-1]) + "."
            count = 1

    # Add the final count and character to the encoded string
    encoded_string += str(count) + "," + str(input_string[-1])

    return encoded_string

In [8]: block_size = 4
#quantization matrix:
quantization_matrix8x8 = np.array([[16, 11, 10, 16, 24, 40, 51, 61],
                                   [12, 12, 14, 19, 26, 58, 60, 55],
                                   [14, 13, 16, 24, 40, 57, 69, 56],
                                   [14, 17, 22, 29, 51, 87, 86, 62],
                                   [18, 22, 37, 56, 68, 109, 103, 77],
                                   [24, 35, 55, 64, 81, 104, 113, 92],
                                   [49, 64, 78, 87, 103, 121, 120, 181],
                                   [72, 92, 95, 98, 112, 100, 103, 99]])
quantization_matrix = np.array([[1,2,4,8],[2,2,4,8],[4,4,4,8],[8,8,8,8]])

In [9]: #calculation for number of blocks needed
h,w = xray.shape
height = np.int32(h)
width = np.int32(w)
h,w = np.float32(h),np.float32(w)

nbh = np.int32(math.ceil(h/block_size))
nbw = np.int32(math.ceil(w/block_size))


#Height of padded image
H = block_size * nbh
W = block_size * nbw

padded_img = np.zeros((H,W))
padded_img[0:height,0:width] = xray[0:height,0:width]

padded_img_minus128 = padded_img - 128
#cv2.imwrite('uncompressed.bmp',np.uint8(padded_img))

In [10]: plt.imshow(padded_img_minus128, cmap='gray')
plt.axis('off')

Out[10]: (-0.5, 1023.5, 1023.5, -0.5)



In [11]: final = np.zeros(xray.shape)
for i in range(nbh):
    #compute start and end row index of block
    row_ind_1 = i*block_size
    row_ind_2 = row_ind_1 + block_size

    for j in range(nbw):
        #compute start and end column index of block
        col_ind_1 = j*block_size
        col_ind_2 = col_ind_1 + block_size

        block = padded_img_minus128[row_ind_1 : row_ind_2, col_ind_1:col_ind_2]

        #applying 2D DCT to block
        DCT = cv2.dct(block)

        quantized = np.round(np.divide(DCT,quantization_matrix))

        reordered = zigzag(quantized)

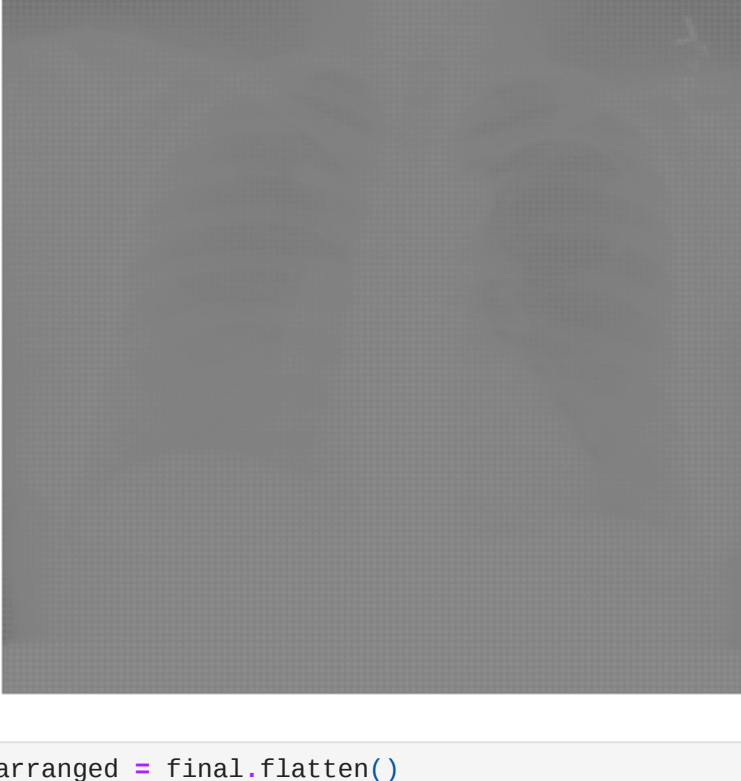
        #reshaped = np.reshape(reordered,(block_size,block_size))

        final[row_ind_1:row_ind_2,col_ind_1:col_ind_2] = reordered

In [12]: plt.imshow(final,cmap='gray')
plt.axis('off')
plt.title("encoded image")
final

Out[12]: array([[[-408., -30., -6., ..., 3., -2., -1],
 [ 4., -7., 1., ..., -2., 1., 0],
 [ 2., -1., 1., ..., -1., -0., 0],
 ...,
 [-25., -29., 5., ..., -29., 8., -4],
 [ 11., 5., -3., ..., 7., 1., 3],
 [ 1., 0., -1., ..., -1., -1., 0]])

        encoded image



In [13]: arranged = final.flatten()

In [14]: arranged.shape

Out[14]: (1048576,)

In [15]: bitstream = run_length_encode(arranged)
#bitstream

In [16]: bitstream2 = str(padded_img.shape[0]) + "," + str(padded_img.shape[1]) + "," + bitstream + ";"

In [17]: file1 = open('test2.txt','w')
file1.write(bitstream2)
file1.close()

In [18]: def r1_decode(g,shape):
    l = g.split(",")
    final = []
    for i in l:
        temp = i.split('.')
        final.append(np.int32(temp))

    #final
    final_arr = []
    for i in final:
        for j in range(i[0]):
            final_arr.append(i[1])

    k = np.array(final_arr)

    return k.reshape(shape)

In [19]: decoded = r1_decode(bitstream,xray.shape)
decoded

Out[19]: array([[[-408., -30., -6., ..., 3., -2., -1],
 [ 4., -7., 1., ..., -2., 1., 0],
 [ 2., -1., 1., ..., -1., 0., 0],
 ...,
 [-25., -29., 5., ..., -29., 8., -4],
 [ 11., 5., -3., ..., 7., 1., 3],
 [ 1., 0., -1., ..., -1., -1., 0]])

In [20]: decoded.shape

Out[20]: (1024, 1024)

In [21]: decoded,final #should be equal to final

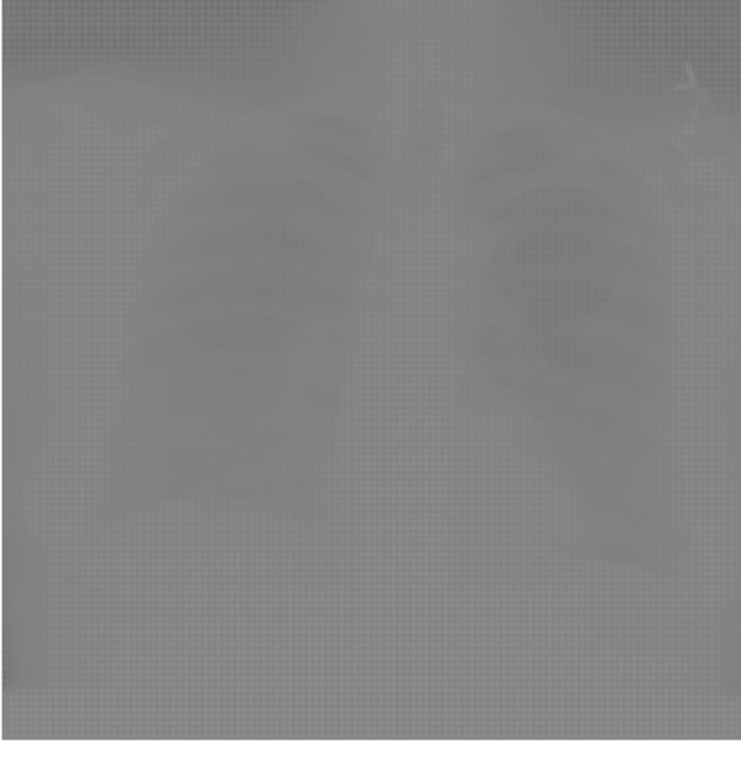
Out[21]: (array([[[-408., -30., -6., ..., 3., -2., -1],
 [ 4., -7., 1., ..., -2., 1., 0],
 [ 2., -1., 1., ..., -1., 0., 0],
 ...,
 [-25., -29., 5., ..., -29., 8., -4],
 [ 11., 5., -3., ..., 7., 1., 3],
 [ 1., 0., -1., ..., -1., -1., 0]]],
 array([[[-408., -30., -6., ..., 3., -2., -1],
 [ 4., -7., 1., ..., -2., 1., 0],
 [ 2., -1., 1., ..., -1., -0., 0],
 ...,
 [-25., -29., 5., ..., -29., 8., -4],
 [ 11., 5., -3., ..., 7., 1., 3],
 [ 1., 0., -1., ..., -1., -1., 0]])])

In [22]: (decoded==final).all()

Out[22]: True

In [23]: plt.imshow(decoded, cmap='gray')
plt.axis('off')

Out[23]: (-0.5, 1023.5, 1023.5, -0.5)



In [24]: uncompressed = np.zeros(decoded.shape)
for i in range(nbh):
    #compute start and end row index of block
    row_ind_1 = i*block_size
    row_ind_2 = row_ind_1 + block_size

    for j in range(nbw):
        #compute start and end column index of block
        col_ind_1 = j*block_size
        col_ind_2 = col_ind_1 + block_size

        block = decoded[row_ind_1 : row_ind_2, col_ind_1:col_ind_2]

        #applying 2D IDCT to block
        reordered = reverse_zigzag_for8x8(block)

        dequantized_block = np.round(np.multiply(block,quantization_matrix)).astype(np.float64)
        IDCT = cv2.idct(dequantized_block)

        #reshaped = np.reshape(reordered,(block_size,block_size))

        uncompressed[row_ind_1:row_ind_2,col_ind_1:col_ind_2] = IDCT

In [25]: uncompressed2 =np.uint8(uncompressed) + 128
uncompressed2

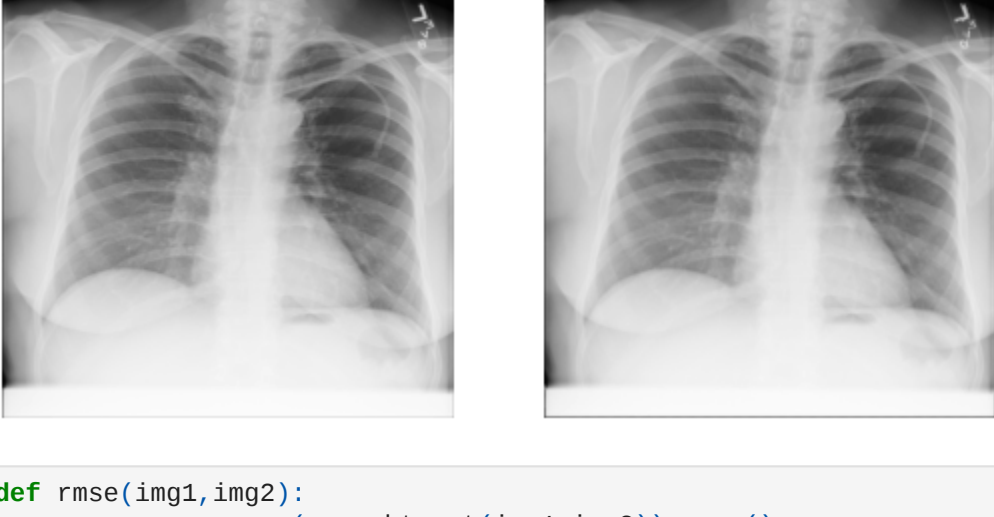
Out[25]: array([[ 0, 25, 44, ..., 15, 11, 14],
 [251, 32, 32, ..., 17, 10, 9],
 [251, 36, 26, ..., 19, 11, 7],
 ...,
 [ 11, 74, 67, ..., 227, 55, 8],
 [ 44, 101, 99, ..., 235, 74, 231],
 [103, 123, 144, ..., 210, 65, 166]], dtype=uint8)

In [26]: (xray==uncompressed2).all()

Out[26]: False

In [27]: plt.subplot(1,2,1)
plt.imshow(xray, cmap='gray')
plt.axis('off')
plt.subplot(1,2,2)
plt.imshow(uncompressed2, cmap='gray')
plt.axis('off')

Out[27]: (-0.5, 1023.5, 1023.5, -0.5)



In [28]: def rmse(img1,img2):
    mse = np.square(np.subtract(img1,img2)).mean()
    rmse = math.sqrt(mse)
    return rmse

def snr(filtered_img,original_img):
    num = np.square(filtered_img).sum()
    diff_sq = np.square(filtered_img - original_img)
    denom = np.sum(diff_sq)
    snr = num/denom
    return snr

In [29]: rmse_jpeg = rmse(uncompressed2, xray)
snr_jpeg = snr(uncompressed2,xray)
print(f'For JPEG: RMSE = {rmse_jpeg} | SNR = {snr_jpeg}')

For JPEG: RMSE = 3.782245347356395 | SNR = 7.51966846572177

In [30]: cv2.imwrite("JPEG Reconstructed Img 4x4 block_size.jpeg",uncompressed2)

Out[30]: True

In [31]: compression_ratio = 117020/1049654
compression_rate = 1/compression_ratio
compression_ratio, compression_rate

Out[31]: (0.11148435579724367, 8.969868398564348)

In [ ]:
```



```
import os

In [18]: import matplotlib.pyplot as plt

In [3]: def walk_through_dir(dir_path):
    """
    Walks through dir_path returning its contents.
    Args:
        dir_path (str or pathlib.Path): target directory

    Returns:
        A print out of:
        - number of subdirectories in dir_path
        - number of images (files) in each subdirectory
        - name of each subdirectory
    """
    for dirpath, dirnames, filenames in os.walk(dir_path):
        print(f"There are {len(dirnames)} directories and {len(filenames)} images in {dirpath}")

In [4]: import random
import zipfile
from pathlib import Path

In [27]: data_path = Path("data/X_ray_image_and_mask_for_U_net_training")
image_path = data_path / "CXR_png"
mask_path = data_path / "mask"

In [6]: walk_through_dir(image_path)

There are 2 directories and 8 images in data\X_ray_image_and_mask_for_U_net_training\CXR_png
There are 0 directories and 288 images in data\X_ray_image_and_mask_for_U_net_training\CXR_png\test
There are 0 directories and 660 images in data\X_ray_image_and_mask_for_U_net_training\CXR_png\train

In [7]: #setup train and testing paths
train_dir = image_path / "train"
test_dir = image_path / "test"

train_dir, test_dir

Out[7]: (WindowsPath('data/X_ray_image_and_mask_for_U_net_training\CXR_png\train'),
WindowsPath('data/X_ray_image_and_mask_for_U_net_training\CXR_png\test'))

In [18]: import random
from PIL import Image

In [68]: #set seed
random.seed(42)

#Get all image paths (" means 'any combination')
image_path_list = list(image_path.glob("**/*.png"))

#Get random image path
random_image_path = random.choice(image_path_list)

#Get image class from path name (name of directory)
image_class = random_image_path.parent.stem

#Open image
img = Image.open(random_image_path)

#print metadata
print(f"Random Image Path: {random_image_path}")
print(f"Image class: {image_class}")
print(f"Image height: {img.height}")
print(f"Image width: {img.width}")

-----
IndexError                                Traceback (most recent call last)
Cell In[68], line 8
      5 image_path_list = list(image_path.glob("**/*.png"))
      7 #Get random image path
----> 8 random_image_path = random.choice(image_path_list)
      9 #Get image class from path name (name of directory)
    11 image_class = random_image_path.parent.stem

File C:\Program Files\WindowsApps\PythonSoftwareFoundation.Python.3.11.3.11.1000.0_x64__qbz5q2kfraagp@Lib\random.py:373, in Random.choice(self, seq)
    378     # As an accommodation for NumPy, we don't use "if not seq"
    379     # because NumPy.array([]) raises a ValueError.
    372 if not len(seq):
--> 373     raise IndexError("Cannot choose from an empty sequence")
    374 return seq[self._randbelow(len(seq))]

IndexError: Cannot choose from an empty sequence

In [16]: #transforming data into tensors, then into torch.utils.data.Dataset
import torch
from torch.utils.data import DataLoader
from torchvision import datasets, transforms

In [124]: data_transform = transforms.Compose([transforms.ToPILImage(),
                                          transforms.Resize(224),
                                          transforms.CenterCrop(224),
                                          transforms.ToTensor()]) #this also converts all pixel values from 0 to 255 to be between 0.0 and 1.0

In [22]: def plot_transformed_images(image_paths, transform, n=3, seed=42):
    """Plots a series of random images from image_paths.

    Will open n image from image_paths, transform them with transform and
    plot them side by side

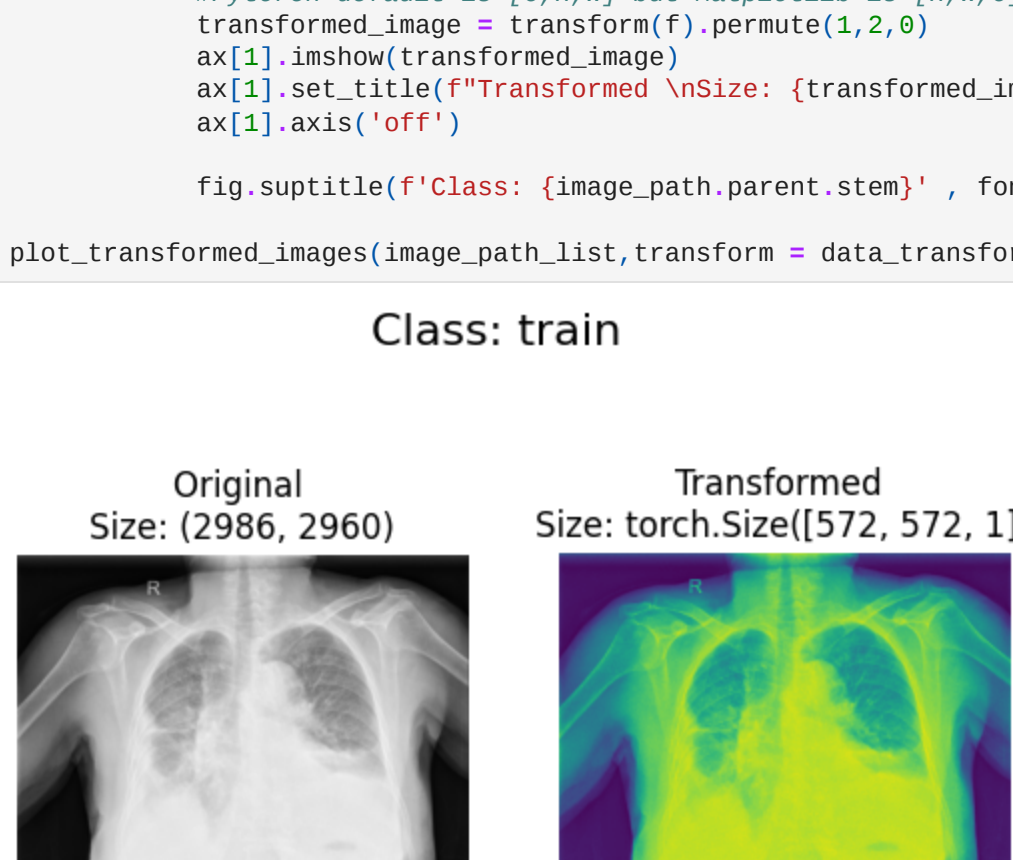
    Args:
        image_path (list): List of target image paths.
        transform (PyTorch Transform): Transforms to apply to images.
        n (int, optional): Number of images to plot. Defaults to 3.
        seed (int, optional): Random seed for the random generator. Defaults to 42
    """
    random.seed(seed)
    random_image_paths = random.sample(image_paths, n)
    for image_path in random_image_paths:
        with Image.open(image_path) as f:
            fig, axs = plt.subplots(1, 2)
            ax[0].imshow(f)
            ax[0].set_title(f"Original \nSize: {f.size}")
            ax[1].axis('off')

            #transform and plot the image
            #note: permute changes shape of image to suit matplotlib
            #pytorch default is [C,H,W] but matplotlib is [H,W,C]
            transformed_image = transform(f).permute(1,2,0)
            ax[1].imshow(transformed_image)
            ax[1].set_title(f"Transformed \nsize: {transformed_image.shape}")
            ax[1].axis('off')


            fig.suptitle(f"Class: {image_path.parent.stem}" , fontsize = 16)

    plot_transformed_images(image_path_list, transform = data_transform, n=3)
```


Class: train



Class: test



Class: test



```
In [24]: #using ImageFolder to create datasets
from torchvision import datasets
train_data = datasets.ImageFolder(root = train_dir,
                                transform = data_transform,
                                target_transform = None)
test_data = datasets.ImageFolder(root = test_dir,
                                transform = data_transform,
                                target_transform = None)
print(f"train data:\n{train_data}\ntest data:\n{test_data}")

train_data:
Dataset ImageFolder
  Number of datapoints: 660
  Root location: data\X_ray_image_and_mask_for_U_net_training\CXR_png\train
  StandardTransform
Transform: Compose(
  Resize(size=(572, 572), interpolation=nearest, max_size=None, antialias=True)
  ToTensor()
)
Test data:
Dataset ImageFolder
  Number of datapoints: 288
  Root location: data\X_ray_image_and_mask_for_U_net_training\CXR_png\test
  StandardTransform
Transform: Compose(
  Resize(size=(572, 572), interpolation=nearest, max_size=None, antialias=True)
  ToTensor()
)
```

```
In [28]: train_data.classes
Out[28]: ['xray']

In [26]: train_data.class_to_idx
Out[26]: {'xray': 0}
```

```
In [135]: #defining U-Net parameters
NUM_CHANNELS = 1
NUM_CLASSES = 1
NUM_LEVELS = 3

#Initialize learning rate, number of epochs to train for and batch size
INIT_LR = 0.001
NUM_EPOCHS = 40
BATCH_SIZE = 64

#define the input image dimensions
INPUT_IMAGE_WIDTH = 320
INPUT_IMAGE_HEIGHT = 128

#threshold to filter weak predictions
THRESHOLD = 0.5

#define path to the base output directory
BASE_OUTPUT = "output"

MODEL_PATH = os.path.join(BASE_OUTPUT, "unet_xray.pth")
PLOT_PATH = os.path.sep.join([BASE_OUTPUT, "plot.png"])
TEST_PATH = os.path.sep.join([BASE_OUTPUT, "test_paths.txt"])
```

```
In [116]: from torch.utils.data import Dataset
class SegmentationDataset(Dataset):
    def __init__(self, imagePaths, maskPaths, transforms):
        self.imagePaths = imagePaths
        self.maskPaths = maskPaths
        self.transforms = transforms

    def __len__(self):
        return len(self.imagePaths)

    def __getitem__(self, idx):
        imagePath = self.imagePaths[idx]
        image = cv2.imread(imagePath)
        image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
        mask = cv2.imread(self.maskPaths[idx], 0)

        # If self.transforms is not None:
        image = self.transforms(image)
        mask = self.transforms(mask)

        return (image, mask)
```

```
In [33]: from torch.nn import ConvTranspose2d
from torch.nn import Conv2d
from torch.nn import MaxPool2d
from torch.nn import Module
from torch.nn import ModuleList
from torch.nn import ReLU
from torchvision.transforms import CenterCrop
from torch.nn import functional as F
```

```
In [98]: import cv2

In [117]: class Block(Module):
    def __init__(self, inChannels, outChannels):
        super().__init__()
        self.conv1 = Conv2d(inChannels, outChannels, 3)
        self.relu = ReLU()
        self.conv2 = Conv2d(outChannels, outChannels, 3)

    def forward(self, x):
        return self.conv2(self.relu(self.conv1(x)))
```

```
In [118]: class Encoder(Module):
    def __init__(self, channels = (3, 16, 32, 64)):
        super().__init__()
        self.pool = MaxPool2d(2)
        self.channels = ModuleList([Block(channels[i], channels[i+1]) for i in range(len(channels)-1)])
        self.pool = MaxPool2d(2)

    def forward(self, x):
        blockOutputs = []
        for block in self.channels:
            x = block(x)
            blockOutputs.append(x)
            x = self.pool(x)
        return blockOutputs
```

```
In [119]: class Decoder(Module):
    def __init__(self, channels=(64, 32, 16)):
        super().__init__()
        # decoder blocks
        self.channels = ModuleList([ConvTranspose2d(channels[i], channels[i+1], 2, 2) for i in range(len(channels)-1)])
        self.dec_blocks = ModuleList([Block(channels[i], channels[i+1]) for i in range(len(channels)-1)])

    def forward(self, x, encFeatures):
        # loop through the number of channels
        for i in range(len(self.channels)-1):
            # pass the inputs through the upsampler blocks
            x = self.upsample(x)
            # crop the current features from the encoder blocks,
            # concatenate them with the current upsampled features,
            # and pass the concatenated output through the current
            # decoder block
            encFeat = self.crop(encFeatures[i], x)
            x = torch.cat([x, encFeat], dim=1)
            x = self.dec_blocks[i](x)
            # return the final decoder output
            return x

    def crop(self, encFeatures, x):
        # grab the dimensions of the inputs, and crop the encoder
        # features to match the dimensions
        [..., H, W] = x.shape
        encFeatures = CenterCrop([H, W])(encFeatures)
        # return the cropped features
        return encFeatures
```

```
In [90]: class UNet(Module):
    def __init__(self, encChannels = (3, 16, 32, 64),
                 decChannels = (64, 32, 16),
                 numClasses = 1,
                 outSize = (INPUT_IMAGE_HEIGHT, INPUT_IMAGE_WIDTH)):
        super().__init__()
        self.encoder = Encoder(encChannels)
        self.decoder = Decoder(decChannels)
        self.head = Conv2d(decChannels[-1], numClasses, 1)
        self.relu = ReLU()
        self.outSize = outSize

    def forward(self, x):
        encFeatures = self.encoder(x)
        encFeatures = self.decoder(encFeatures[:-1][0])
        map = self.head(decFeatures)

        if self.relu:
            map = F.interpolate(map, self.outSize)

        return map
```

```
In [84]: from torch.nn import BCEWithLogitsLoss
from torch.optim import Adam
from torch.utils.data import DataLoader
from sklearn.model_selection import train_test_split
from torchvision import transforms
from torch.nn import functional as F
import matplotlib.pyplot as plt
import time
import os
from torch import nn
```

```
In [120]: imagePaths = sorted(list(paths_list_images(image_path)))
maskPaths = sorted(list(paths_list_images(mask_path)))

split = train_test_split(imagePaths, maskPaths, test_size = 0.25, random_state = 42)
(trainImages, testImages) = split[:2]
(trainMasks, testMasks) = split[2:]

print(f"[INFO] saving testing image paths")
f = open(test_dir, "w")
f.write("\n".join(testImages))
f.close()

[INFO] saving testing image paths
```

```
In [125]: trainDS = SegmentationDataset(imagePaths = trainImages,
                                   maskPaths = trainMasks,
                                   transforms = data_transform)
testDS = SegmentationDataset(imagePaths = testImages,
                             maskPaths = testMasks,
                             transforms = data_transform)

trainLoader = DataLoader(trainDS, shuffle = True,
                        batch_size = BATCH_SIZE)
testLoader = DataLoader(testDS, shuffle = False, batch_size = BATCH_SIZE)
```

```
In [126]: len(trainDS[0][0])
Out[126]: 3

In [127]: trainDS[0][0].size()
Out[127]: torch.Size([3, 128, 128])
```

```
In [130]: unet = UNet()
lossfn = BCEWithLogitsLoss()
optimizer = Adam(unet.parameters()) lr = INIT_LR
trainSteps = len(trainDS) // BATCH_SIZE
testSteps = len(testDS) // BATCH_SIZE
H = ("train_loss" : [], "test_loss" : [])
```

```
In [132]: startTime = time.time()
for e in tqdm(range(NUM_EPOCHS)):
    unet.train()
    totalTrainLoss = 0
    totalTestLoss = 0

    for i, (x, y) in enumerate(trainLoader):
        pred = unet(x)
        loss = lossfn(pred, y)
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        totalTrainLoss += loss

    with torch.inference_mode():
        unet.eval()

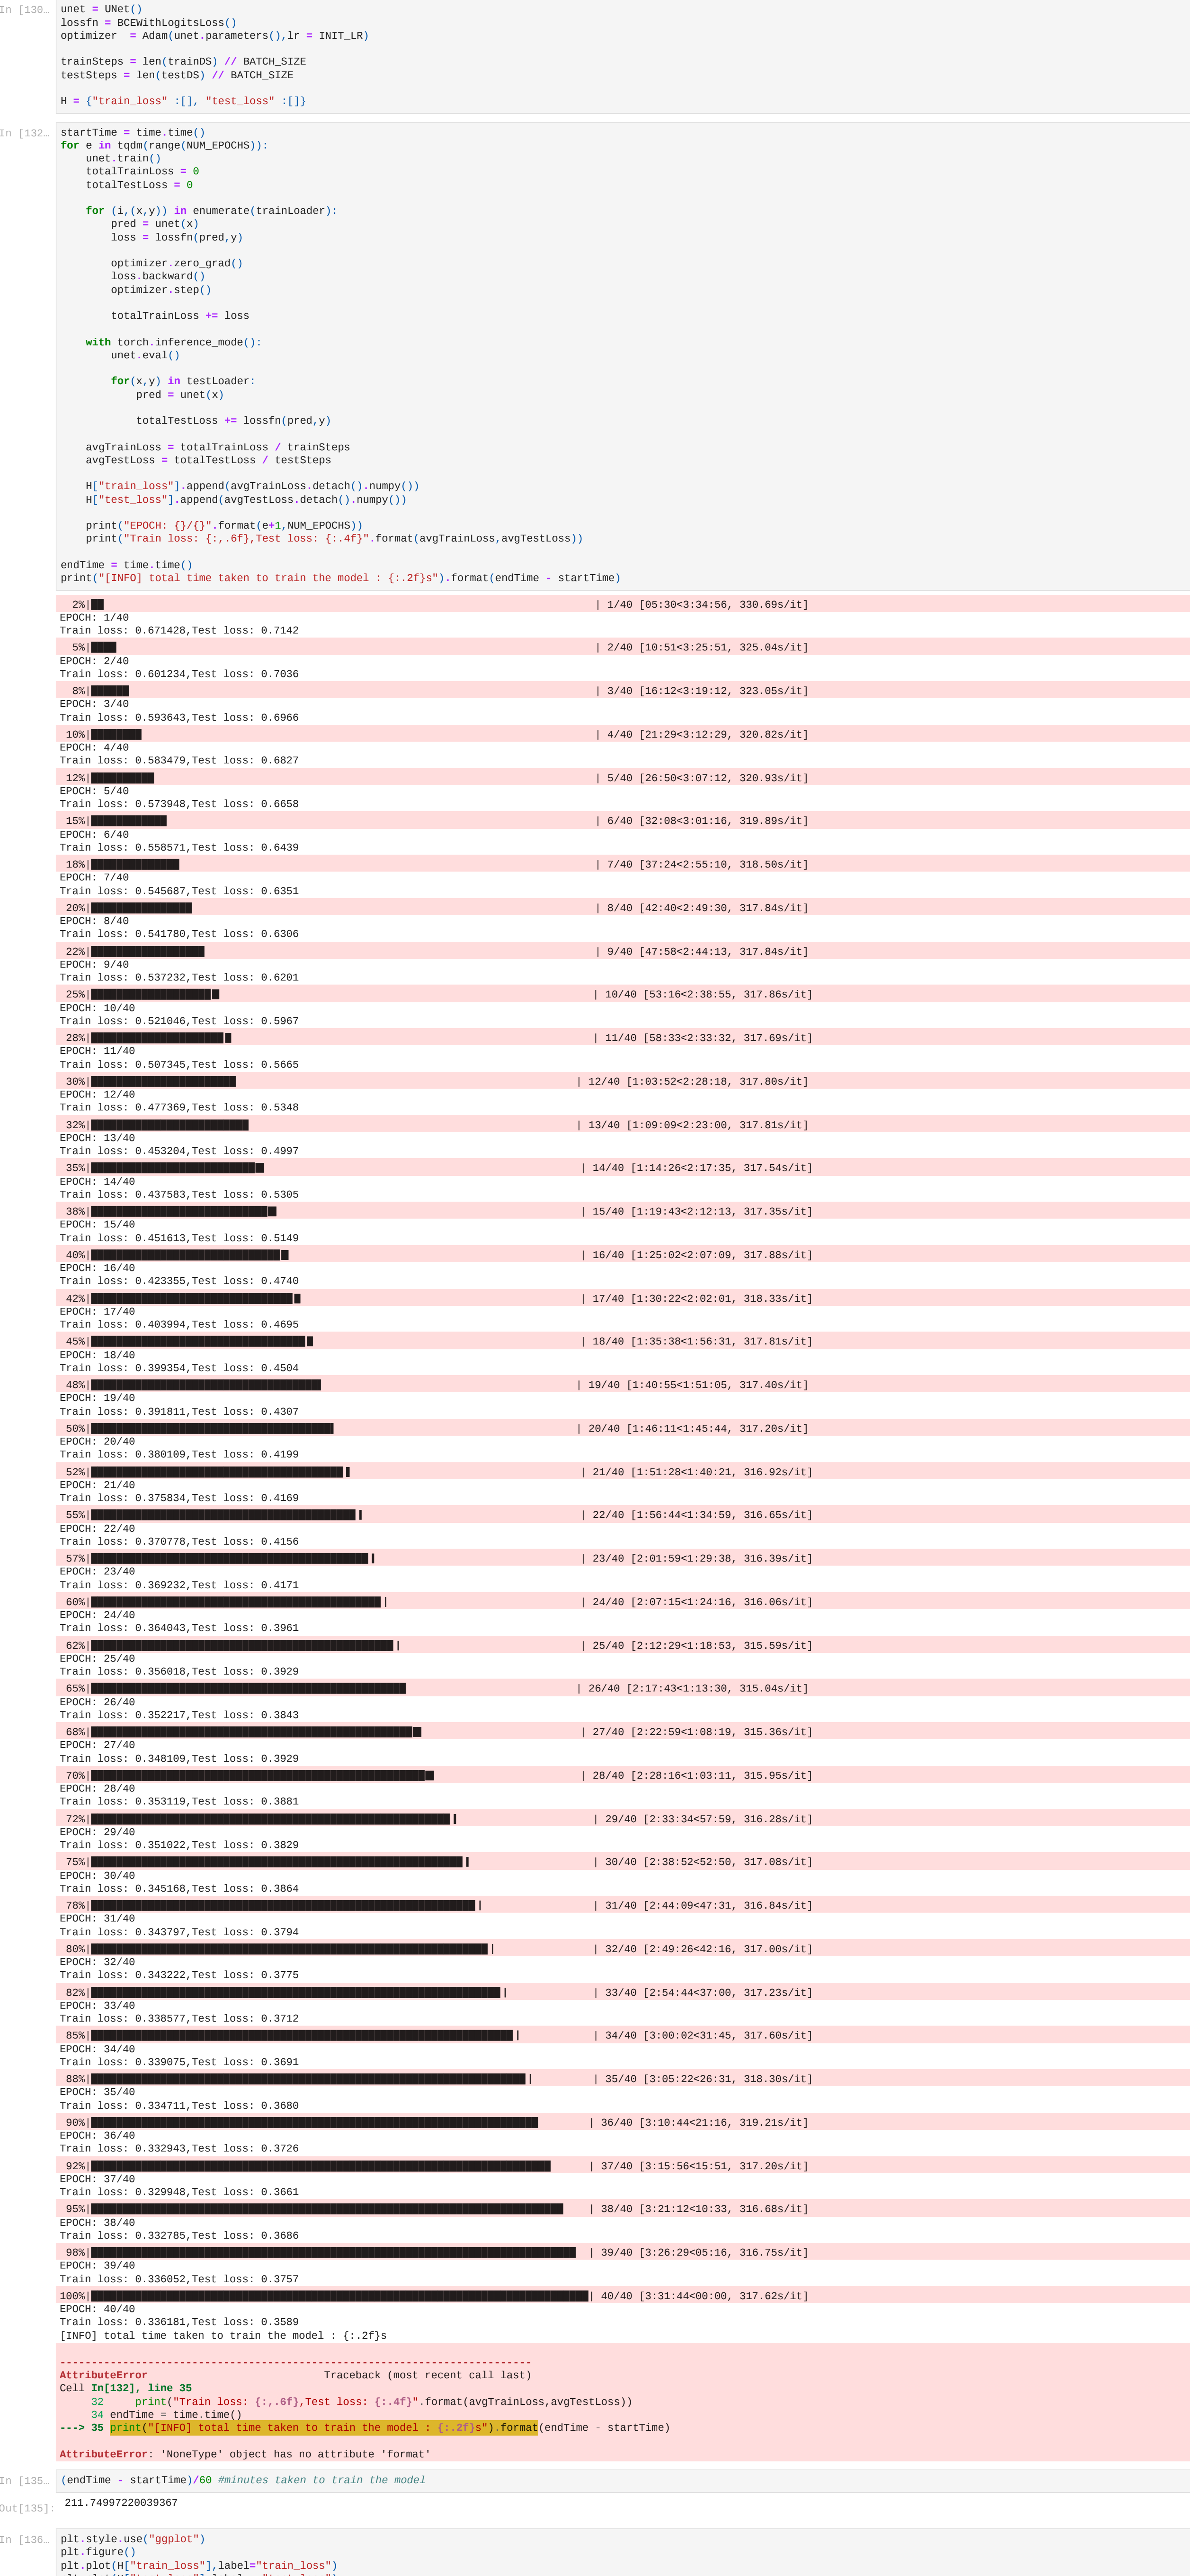
        for (x, y) in testLoader:
            pred = unet(x)
            totalTestLoss += lossfn(pred, y)

    avgTrainLoss = totalTrainLoss / trainSteps
    avgTestLoss = totalTestLoss / testSteps

    H["train_loss"].append(avgTrainLoss.detach().numpy())
    H["test_loss"].append(avgTestLoss.detach().numpy())

    print(f"EPOCH: {i+1} | {f'{avgTrainLoss:.4f}'} | {f'{avgTestLoss:.4f}'}")
    print(f"Train loss: {f'{avgTrainLoss:.4f}'} | Test loss: {f'{avgTestLoss:.4f}'}")

endTime = time.time()
print(f"[INFO] total time taken to train the model : {f'{(endTime - startTime):.2f}'}")
```



```
In [138]: torch.save(unet, MODEL_PATH)

Making Predictions

In [147]: import numpy as np

In [156]: def prepare_plot(origImage, predMask):
    figure, ax = plt.subplots(nrows = 1, ncols = 2, figsize = (10,10))

    ax[0].imshow(origImage)
    ax[1].imshow(predMask)
    ax[1].imshow(predMask, cmap='gray')

    ax[0].set_title("Image")
    ax[1].set_title("Original Mask")
    ax[1].set_title("Predicted Mask")

    figure.tight_layout()
    figure.show()

In [157]: unet.eval()

with torch.inference_mode():
    image = cv2.imread('data/X_ray_image_and_mask_for_U_net_training/CXR_png/CNCRX_0035_0.png')
    image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
    image = image.astype("float32")/255.0
    image = torch.tensor(image, [2, 255, 0])

    orig = image.copy()

    # filename = the path, split on path.sep[-1]
    groundTruthPath = os.path.join(mask_path, filename)

    gtMask = cv2.imread(groundTruthPath, 0)
    gtMask = cv2.resize(gtMask, (INPUT_IMAGE_HEIGHT, INPUT_IMAGE_WIDTH))

    image = np.transpose(image, (2, 0, 1)) #to get it to [C, H, W] format
    image = np.expand_dims(image, 0) #model inputs 4 dimensional tensor
    image = torch.tensor(image)

    predMask = unet(image).squeeze() #to remove extra dimension
    predMask = torch.sigmoid(predMask)
    predMask = predMask.numpy()

    prepare_plot(orig, predMask)
```

C:\Users\ayay\AppData\Local\Temp\ipynbkernel-42868\4862992238.py:13: UserWarning: Matplotlib is currently using module://matplotlib_inline.backend_inline, which is a non-GUI backend, so cannot show the Figure.



```
In [138]: torch.save(unet, MODEL_PATH)

Making Predictions

In [147]: import numpy as np

In [156]: def prepare_plot(origImage, predMask):
    figure, ax = plt.subplots(nrows = 1, ncols = 2, figsize = (10,10))

    ax[0].imshow(origImage)
    ax[1].imshow(predMask)
    ax[1].imshow(predMask, cmap='gray')

    ax[0].set_title("Image")
    ax[1].set_title("Original Mask")
    ax[1].set_title("Predicted Mask")

    figure.tight_layout()
    figure.show()

In [157]: unet.eval()

with torch.inference_mode():
    image = cv2.imread('data/X_ray_image_and_mask_for_U_net_training/CXR_png/CNCRX_0035_0.png')
    image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
    image = image.astype("float32")/255.0
    image = torch.tensor(image, [2, 255, 0])

    orig = image.copy()

    # filename = the path, split on path.sep[-1]
    groundTruthPath = os.path.join(mask_path, filename)

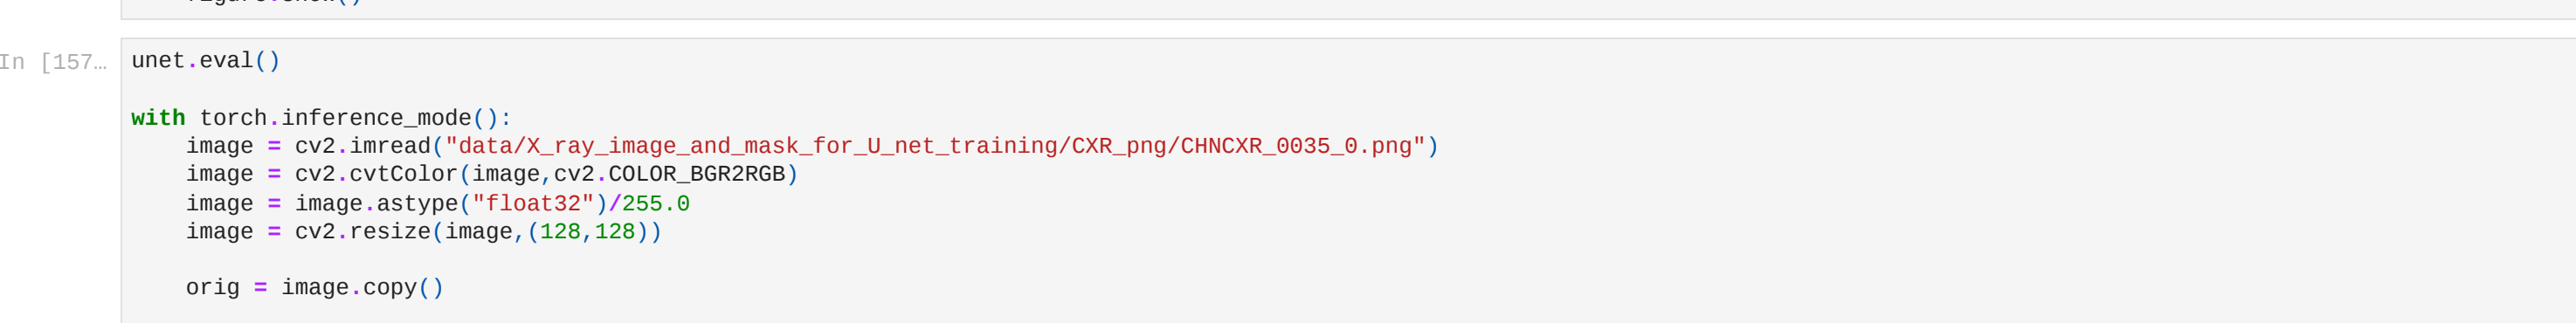
    gtMask = cv2.imread(groundTruthPath, 0)
    gtMask = cv2.resize(gtMask, (INPUT_IMAGE_HEIGHT, INPUT_IMAGE_WIDTH))

    image = np.transpose(image, (2, 0, 1)) #to get it to [C, H, W] format
    image = np.expand_dims(image, 0) #model inputs 4 dimensional tensor
    image = torch.tensor(image)

    predMask = unet(image).squeeze() #to remove extra dimension
    predMask = torch.sigmoid(predMask)
    predMask = predMask.numpy()

    prepare_plot(orig, predMask)
```

C:\Users\ayay\AppData\Local\Temp\ipynbkernel-42868\4862992238.py:13: UserWarning: Matplotlib is currently using module://matplotlib_inline.backend_inline, which is a non-GUI backend, so cannot show the Figure.




```
In [1]: # Importing necessary libraries
import os
import random
import shutil
import numpy as np
import matplotlib.pyplot as plt
from keras.preprocessing.image import ImageDataGenerator
from keras.applications import VGG16, VGG19, ResNet50, InceptionV3, Xception
from keras.models import Model
from keras.layers import Dense, GlobalAveragePooling2D
from keras.optimizers import Adam
from keras.callbacks import ModelCheckpoint
from sklearn.metrics import classification_report, confusion_matrix

# Setting random seed for reproducibility
np.random.seed(42)
random.seed(42)

# Define constants
NUM_CLASSES = 2
IMG_SIZE = (224, 224)
BATCH_SIZE = 32
TRAIN_DIR = 'train'
VAL_DIR = 'val'
TEST_DIR = 'test'
TRAIN_SPLIT = 0.7
VAL_SPLIT = 0.1
TEST_SPLIT = 0.2
test_samples = 280

In [2]: # Define the base directory where the image data is located
BASE_DIR = r"C:\Users\mohit\Downloads\DIP Assignment 2\Data"

In [3]: # Set the paths for training, validation, and test data
TRAIN_DIR = os.path.join(BASE_DIR, 'train')
VAL_DIR = os.path.join(BASE_DIR, 'val')
TEST_DIR = os.path.join(BASE_DIR, 'test')

In [4]: # Load and preprocess the data
train_datagen = ImageDataGenerator(rescale=1./255, shear_range=0.2, zoom_range=0.2, horizontal_flip=True)
val_datagen = ImageDataGenerator(rescale=1./255)
test_datagen = ImageDataGenerator(rescale=1./255)

In [5]: train_generator = train_datagen.flow_from_directory(TRAIN_DIR, target_size=IMG_SIZE, batch_size=BATCH_SIZE, class_mode='categorical')

val_generator = val_datagen.flow_from_directory(VAL_DIR, target_size=IMG_SIZE, batch_size=BATCH_SIZE, class_mode='categorical')

test_generator = test_datagen.flow_from_directory(TEST_DIR, target_size=IMG_SIZE, batch_size=BATCH_SIZE, class_mode='categorical', shuffle=False)

Found 980 images belonging to 2 classes.
Found 140 images belonging to 2 classes.
Found 280 images belonging to 2 classes.

In [6]: # Define function for creating transfer learning models
def create_transfer_model(base_model, num_classes):
    x = base_model.output
    x = GlobalAveragePooling2D()(x)
    x = Dense(1024, activation='relu')(x)
    predictions = Dense(num_classes, activation='softmax')(x)
    model = Model(inputs=base_model.input, outputs=predictions)
    return model

# Define VGG16 model
vgg16_base = VGG16(include_top=False, weights='imagenet', input_shape=(IMG_SIZE[0], IMG_SIZE[1], 3))
vgg16_model = create_transfer_model(vgg16_base, NUM_CLASSES)

# Define VGG19 model
vgg19_base = VGG19(include_top=False, weights='imagenet', input_shape=(IMG_SIZE[0], IMG_SIZE[1], 3))
vgg19_model = create_transfer_model(vgg19_base, NUM_CLASSES)

In [7]: # Compile the models
optimizer = Adam(learning_rate=0.0001)
vgg16_model.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=['accuracy'])
vgg19_model.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=['accuracy'])

In [8]: # Define checkpoint callback to save the best model during training
checkpoint = ModelCheckpoint('best_model.h5', monitor='val_loss', save_best_only=True)

In [9]: # Train the models
vgg16_history = vgg16_model.fit_generator(train_generator, steps_per_epoch=train_generator.n // train_generator.batch_size, epochs=5, validation_data=val_generator, validation_steps=val_generator.n // val_generator.batch_size, callbacks=[checkpoint])
vgg19_history = vgg19_model.fit_generator(train_generator, steps_per_epoch=train_generator.n // train_generator.batch_size, epochs=5, validation_data=val_generator, validation_steps=val_generator.n // val_generator.batch_size, callbacks=[checkpoint])

C:\Users\mohit\AppData\Local\Temp\ipykernel_25572\2039408895.py:2: UserWarning: `Model.fit_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.
  vgg16_history = vgg16_model.fit_generator(train_generator, steps_per_epoch=train_generator.n // train_generator.batch_size, epochs=5, validation_data=val_generator, validation_steps=val_generator.n // val_generator.batch_size, callbacks=[checkpoint])
Epoch 1/5
30/30 [=====] - 269s 9s/step - loss: 0.6998 - accuracy: 0.5190 - val_loss: 0.7125 - val_accuracy: 0.5000
Epoch 2/5
30/30 [=====] - 286s 10s/step - loss: 0.6402 - accuracy: 0.6118 - val_loss: 0.4709 - val_accuracy: 0.8281
Epoch 3/5
30/30 [=====] - 470s 16s/step - loss: 0.2767 - accuracy: 0.9008 - val_loss: 0.1435 - val_accuracy: 0.9531
Epoch 4/5
30/30 [=====] - 503s 17s/step - loss: 0.1493 - accuracy: 0.9473 - val_loss: 0.1466 - val_accuracy: 0.9609
Epoch 5/5
30/30 [=====] - 504s 17s/step - loss: 0.1644 - accuracy: 0.9494 - val_loss: 0.0989 - val_accuracy: 0.9688
C:\Users\mohit\AppData\Local\Temp\ipykernel_25572\2039408895.py:3: UserWarning: `Model.fit_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.
  vgg19_history = vgg19_model.fit_generator(train_generator, steps_per_epoch=train_generator.n // train_generator.batch_size, epochs=5, validation_data=val_generator, validation_steps=val_generator.n // val_generator.batch_size, callbacks=[checkpoint])
Epoch 1/5
30/30 [=====] - 570s 19s/step - loss: 0.7282 - accuracy: 0.5042 - val_loss: 0.6897 - val_accuracy: 0.5156
Epoch 2/5
30/30 [=====] - 558s 19s/step - loss: 0.6453 - accuracy: 0.6983 - val_loss: 0.4191 - val_accuracy: 0.8672
Epoch 3/5
30/30 [=====] - 544s 18s/step - loss: 0.3698 - accuracy: 0.8576 - val_loss: 0.1742 - val_accuracy: 0.9297
Epoch 4/5
30/30 [=====] - 547s 18s/step - loss: 0.2672 - accuracy: 0.9040 - val_loss: 0.1718 - val_accuracy: 0.9375
Epoch 5/5
30/30 [=====] - 354s 12s/step - loss: 0.2497 - accuracy: 0.9146 - val_loss: 0.2170 - val_accuracy: 0.8984

In [10]: # Evaluate the models on test data
vgg16_scores = vgg16_model.evaluate_generator(test_generator, steps=test_samples // BATCH_SIZE)
print("VGG16 Test Loss:", vgg16_scores[0])
print("VGG16 Test Accuracy:", vgg16_scores[1])

vgg19_scores = vgg19_model.evaluate_generator(test_generator, steps=test_samples // BATCH_SIZE)
print("VGG19 Test Loss:", vgg19_scores[0])
print("VGG19 Test Accuracy:", vgg19_scores[1])

C:\Users\mohit\AppData\Local\Temp\ipykernel_25572\1392354301.py:2: UserWarning: `Model.evaluate_generator` is deprecated and will be removed in a future version. Please use `Model.evaluate`, which supports generators.
  vgg16_scores = vgg16_model.evaluate_generator(test_generator, steps=test_samples // BATCH_SIZE)
VGG16 Test Loss: 0.07796119153499603
VGG16 Test Accuracy: 0.97265625
C:\Users\mohit\AppData\Local\Temp\ipykernel_25572\1392354301.py:6: UserWarning: `Model.evaluate_generator` is deprecated and will be removed in a future version. Please use `Model.evaluate`, which supports generators.
  vgg19_scores = vgg19_model.evaluate_generator(test_generator, steps=test_samples // BATCH_SIZE)
VGG19 Test Loss: 0.14393511414527893
VGG19 Test Accuracy: 0.94140625

In [15]: # Get the number of test samples
num_test_samples = test_generator.n

# Calculate the number of steps for prediction
num_prediction_steps = num_test_samples // test_generator.batch_size + 1

# Generate predictions for all test samples
vgg16_predictions = vgg16_model.predict_generator(test_generator, steps=num_prediction_steps, verbose=1)

# Convert predictions to class labels
vgg16_predicted_labels = np.argmax(vgg16_predictions, axis=1)

C:\Users\mohit\AppData\Local\Temp\ipykernel_25572\2218791254.py:8: UserWarning: `Model.predict_generator` is deprecated and will be removed in a future version. Please use `Model.predict`, which supports generators.
  vgg16_predictions = vgg16_model.predict_generator(test_generator, steps=num_prediction_steps, verbose=1)
9/9 [=====] - 23s 3s/step

In [16]: # Get the number of test samples
num_test_samples = test_generator.n

# Calculate the number of steps for prediction
num_prediction_steps = num_test_samples // test_generator.batch_size + 1

# Generate predictions for all test samples
vgg19_predictions = vgg19_model.predict_generator(test_generator, steps=num_prediction_steps, verbose=1)

# Convert predictions to class labels
vgg19_predicted_labels = np.argmax(vgg19_predictions, axis=1)

C:\Users\mohit\AppData\Local\Temp\ipykernel_25572\903632277.py:8: UserWarning: `Model.predict_generator` is deprecated and will be removed in a future version. Please use `Model.predict`, which supports generators.
  vgg19_predictions = vgg19_model.predict_generator(test_generator, steps=num_prediction_steps, verbose=1)
9/9 [=====] - 30s 3s/step

In [17]: # Get true class labels
true_labels = test_generator.classes

In [18]: # Calculate classification report
vgg16_report = classification_report(true_labels, vgg16_predicted_labels)

In [19]: # Calculate classification report
from sklearn.metrics import classification_report

# Get the ground truth labels
ground_truth_labels = test_generator.classes

# Get the predicted labels
vgg19_predicted_labels = np.argmax(vgg19_predictions, axis=1)

# Calculate classification report
classification_report_vgg19 = classification_report(ground_truth_labels, vgg19_predicted_labels, zero_division=1)

In [20]: print("vgg16 Classification Report:")
print(vgg16_report)

print("vgg19 Classification Report:")
print(classification_report_vgg19)

vgg16 Classification Report:
      precision    recall  f1-score   support

     0       0.95      0.99      0.97         140
     1       0.99      0.94      0.96         140

   accuracy
 macro avg       0.97      0.96      0.96         280
weighted avg       0.97      0.96      0.96         280

vgg19 Classification Report:
      precision    recall  f1-score   support

     0       0.87      0.99      0.93         140
     1       0.98      0.86      0.92         140

   accuracy
 macro avg       0.93      0.92      0.92         280
weighted avg       0.93      0.92      0.92         280
```

In []:

```
In [5]: # Importing necessary libraries
import os
import random
import shutil
import numpy as np
import matplotlib.pyplot as plt
from keras.preprocessing.image import ImageDataGenerator
from keras.applications import VGG16, VGG19, ResNet50, InceptionV3, Xception
from keras.models import Model
from keras.layers import Dense, GlobalAveragePooling2D
from keras.optimizers import Adam
from keras.callbacks import ModelCheckpoint
from sklearn.metrics import classification_report, confusion_matrix

# Setting random seed for reproducibility
np.random.seed(42)
random.seed(42)

# Define constants
NUM_CLASSES = 2
IMG_SIZE = (224, 224)
BATCH_SIZE = 32
TRAIN_DIR = 'train'
VAL_DIR = 'val'
TEST_DIR = 'test'
TRAIN_SPLIT = 0.7
VAL_SPLIT = 0.1
TEST_SPLIT = 0.2
test_samples = 280

# Define the base directory where the image data is located
BASE_DIR = r"C:\Users\mohit\Downloads\DIP Assignment 2\Data"

# Set the paths for training, validation, and test data
TRAIN_DIR = os.path.join(BASE_DIR, 'train')
VAL_DIR = os.path.join(BASE_DIR, 'val')
TEST_DIR = os.path.join(BASE_DIR, 'test')

# Load and preprocess the data
train_datagen = ImageDataGenerator(rescale=1./255, shear_range=0.2, zoom_range=0.2, horizontal_flip=True)
val_datagen = ImageDataGenerator(rescale=1./255)
test_datagen = ImageDataGenerator(rescale=1./255)

train_generator = train_datagen.flow_from_directory(TRAIN_DIR, target_size=IMG_SIZE, batch_size=BATCH_SIZE, class_mode='categorical')

val_generator = val_datagen.flow_from_directory(VAL_DIR, target_size=IMG_SIZE, batch_size=BATCH_SIZE, class_mode='categorical')

test_generator = test_datagen.flow_from_directory(TEST_DIR, target_size=IMG_SIZE, batch_size=BATCH_SIZE, class_mode='categorical', shuffle=False)

# Define function for creating transfer learning models
def create_transfer_model(base_model, num_classes):
    x = base_model.output
    x = GlobalAveragePooling2D()(x)
    x = Dense(1024, activation='relu')(x)
    predictions = Dense(num_classes, activation='softmax')(x)
    model = Model(inputs=base_model.input, outputs=predictions)
    return model

# Define Xception model
xception_base = Xception(include_top=False, weights='imagenet', input_shape=(IMG_SIZE[0], IMG_SIZE[1], 3))
xception_model = create_transfer_model(xception_base, NUM_CLASSES)

# Compile the models
optimizer = Adam(learning_rate=0.0001)
xception_model.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=['accuracy'])

# Define checkpoint callback to save the best model during training
checkpoint = ModelCheckpoint('best_model.h5', monitor='val_loss', save_best_only=True)

xception_model.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=['accuracy'])

xception_history = xception_model.fit_generator(train_generator, steps_per_epoch=train_generator.n // train_generator.batch_size, epochs=5, validation_data=val_generator, validation_steps=val_generator.n // val_generator.batch_size, callbacks=[checkpoint])

xception_scores = xception_model.evaluate_generator(test_generator, steps=test_samples // BATCH_SIZE)
print("Xception Test Loss:", xception_scores[0])
print("Xception Test Accuracy:", xception_scores[1])

# Get the number of test samples
num_test_samples = test_generator.n

# Calculate the number of steps for prediction
num_prediction_steps = num_test_samples // test_generator.batch_size + 1

# Generate predictions for all test samples
xception_predictions = xception_model.predict_generator(test_generator, steps=num_prediction_steps, verbose=1)

# Convert predictions to class labels
xception_predicted_labels = np.argmax(xception_predictions, axis=1)

# Get true class labels
true_labels = test_generator.classes

# Calculate classification report
xception_report = classification_report(true_labels, xception_predicted_labels)

print("Xception Classification Report:")
print(xception_report)

Found 980 images belonging to 2 classes.
Found 140 images belonging to 2 classes.
Found 280 images belonging to 2 classes.
C:\Users\mohit\AppData\Local\Temp\ipykernel_33056\3483214989.py:72: UserWarning: `Model.fit_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.
    xception_history = xception_model.fit_generator(train_generator, steps_per_epoch=train_generator.n // train_generator.batch_size, epochs=5, validation_data=val_generator, validation_steps=val_generator.n // val_generator.batch_size, callbacks=[checkpoint])
Epoch 1/5
30/30 [=====] - 186s 6s/step - loss: 0.1842 - accuracy: 0.9462 - val_loss: 0.1311 - val_accuracy: 0.9219
Epoch 2/5
30/30 [=====] - 193s 6s/step - loss: 0.0225 - accuracy: 0.9926 - val_loss: 0.1422 - val_accuracy: 0.9297
Epoch 3/5
30/30 [=====] - 190s 6s/step - loss: 0.0069 - accuracy: 0.9989 - val_loss: 0.0426 - val_accuracy: 0.9922
Epoch 4/5
30/30 [=====] - 192s 6s/step - loss: 0.0070 - accuracy: 0.9979 - val_loss: 0.0151 - val_accuracy: 0.9922
Epoch 5/5
30/30 [=====] - 192s 6s/step - loss: 0.0016 - accuracy: 1.0000 - val_loss: 0.0074 - val_accuracy: 0.9922
C:\Users\mohit\AppData\Local\Temp\ipykernel_33056\3483214989.py:74: UserWarning: `Model.evaluate_generator` is deprecated and will be removed in a future version. Please use `Model.evaluate`, which supports generators.
    xception_scores = xception_model.evaluate_generator(test_generator, steps=test_samples // BATCH_SIZE)
Xception Test Loss: 0.012623521499335766
Xception Test Accuracy: 0.9921875
C:\Users\mohit\AppData\Local\Temp\ipykernel_33056\3483214989.py:85: UserWarning: `Model.predict_generator` is deprecated and will be removed in a future version. Please use `Model.predict`, which supports generators.
    xception_predictions = xception_model.predict_generator(test_generator, steps=num_prediction_steps, verbose=1)
9/9 [=====] - 11s 1s/step
Xception Classification Report:
      precision    recall  f1-score   support

      0         0.99      1.00      0.99        140
      1         1.00      0.99      0.99        140

   accuracy         0.99
  macro avg         0.99      0.99      0.99        280
 weighted avg         0.99      0.99      0.99        280
```



```
In [1]: # Importing necessary libraries
import os
import random
import shutil
import numpy as np
import matplotlib.pyplot as plt
from keras.preprocessing.image import ImageDataGenerator
from keras.applications import VGG16, VGG19, ResNet50, InceptionV3, Xception
from keras.models import Model
from keras.layers import Dense, GlobalAveragePooling2D
from keras.optimizers import Adam
from keras.callbacks import ModelCheckpoint
from sklearn.metrics import classification_report, confusion_matrix

# Setting random seed for reproducibility
np.random.seed(42)
random.seed(42)

# Define constants
NUM_CLASSES = 2
IMG_SIZE = (224, 224)
BATCH_SIZE = 32
TRAIN_DIR = 'train'
VAL_DIR = 'val'
TEST_DIR = 'test'
TRAIN_SPLIT = 0.7
VAL_SPLIT = 0.1
TEST_SPLIT = 0.2
test_samples = 280

In [2]: # Define the base directory where the image data is located
BASE_DIR = r"C:\Users\mohit\Downloads\DIP Assignment 2\Data"

In [3]: # Set the paths for training, validation, and test data
TRAIN_DIR = os.path.join(BASE_DIR, 'train')
VAL_DIR = os.path.join(BASE_DIR, 'val')
TEST_DIR = os.path.join(BASE_DIR, 'test')

In [4]: # Load and preprocess the data
train_datagen = ImageDataGenerator(rescale=1./255, shear_range=0.2, zoom_range=0.2, horizontal_flip=True)
val_datagen = ImageDataGenerator(rescale=1./255)
test_datagen = ImageDataGenerator(rescale=1./255)

In [5]: train_generator = train_datagen.flow_from_directory(TRAIN_DIR, target_size=IMG_SIZE, batch_size=BATCH_SIZE, class_mode='categorical')

val_generator = val_datagen.flow_from_directory(VAL_DIR, target_size=IMG_SIZE, batch_size=BATCH_SIZE, class_mode='categorical')

test_generator = test_datagen.flow_from_directory(TEST_DIR, target_size=IMG_SIZE, batch_size=BATCH_SIZE, class_mode='categorical', shuffle=False)

Found 980 images belonging to 2 classes.
Found 140 images belonging to 2 classes.
Found 280 images belonging to 2 classes.

In [6]: # Define function for creating transfer learning models
def create_transfer_model(base_model, num_classes):
    x = base_model.output
    x = GlobalAveragePooling2D()(x)
    x = Dense(1024, activation='relu')(x)
    predictions = Dense(num_classes, activation='softmax')(x)
    model = Model(inputs=base_model.input, outputs=predictions)
    return model

In [7]: # Define InceptionV3 model
inceptionv3_base = InceptionV3(include_top=False, weights='imagenet', input_shape=(IMG_SIZE[0], IMG_SIZE[1], 3))
inceptionv3_model = create_transfer_model(inceptionv3_base, NUM_CLASSES)

In [8]: # Compile the models
optimizer = Adam(learning_rate=0.0001)
inceptionv3_model.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=['accuracy'])

In [9]: # Define checkpoint callback to save the best model during training
checkpoint = ModelCheckpoint('best_model.h5', monitor='val_loss', save_best_only=True)

In [10]: # Train the models
inceptionv3_history = inceptionv3_model.fit_generator(train_generator, steps_per_epoch=train_generator.n // train_generator.batch_size, epochs=10, validation_data=val_generator, va

C:\Users\mohit\AppData\Local\Temp\ipykernel_34676\918017655.py:2: UserWarning: `Model.fit_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.
    inceptionv3_history = inceptionv3_model.fit_generator(train_generator, steps_per_epoch=train_generator.n // train_generator.batch_size, epochs=10, validation_data=val_generator, validation_steps=val_generator.n // val_generator.batch_size, callbacks=[checkpoint])
Epoch 1/10
30/30 [=====] - 110s 4s/step - loss: 0.1248 - accuracy: 0.9420 - val_loss: 0.0394 - val_accuracy: 0.9922
Epoch 2/10
30/30 [=====] - 112s 4s/step - loss: 0.0094 - accuracy: 0.9979 - val_loss: 0.0310 - val_accuracy: 0.9922
Epoch 3/10
30/30 [=====] - 112s 4s/step - loss: 0.0092 - accuracy: 0.9968 - val_loss: 0.0544 - val_accuracy: 0.9922
Epoch 4/10
30/30 [=====] - 112s 4s/step - loss: 0.0117 - accuracy: 0.9979 - val_loss: 6.0826e-04 - val_accuracy: 1.0000
Epoch 5/10
30/30 [=====] - 111s 4s/step - loss: 0.0068 - accuracy: 0.9979 - val_loss: 0.0132 - val_accuracy: 0.9922
Epoch 6/10
30/30 [=====] - 111s 4s/step - loss: 0.0015 - accuracy: 1.0000 - val_loss: 0.0223 - val_accuracy: 0.9844
Epoch 7/10
30/30 [=====] - 113s 4s/step - loss: 0.0056 - accuracy: 0.9968 - val_loss: 7.1364e-04 - val_accuracy: 1.0000
Epoch 8/10
30/30 [=====] - 113s 4s/step - loss: 0.0047 - accuracy: 0.9968 - val_loss: 0.0220 - val_accuracy: 0.9844
Epoch 9/10
30/30 [=====] - 112s 4s/step - loss: 0.0020 - accuracy: 1.0000 - val_loss: 0.0742 - val_accuracy: 0.9844
Epoch 10/10
30/30 [=====] - 111s 4s/step - loss: 0.0023 - accuracy: 0.9989 - val_loss: 0.0065 - val_accuracy: 1.0000

In [11]: inceptionv3_scores = inceptionv3_model.evaluate_generator(test_generator, steps=test_samples // BATCH_SIZE)
print("inceptionv3 Test Loss:", inceptionv3_scores[0])
print("inceptionv3 Test Accuracy:", inceptionv3_scores[1])

C:\Users\mohit\AppData\Local\Temp\ipykernel_34676\1195483549.py:1: UserWarning: `Model.evaluate_generator` is deprecated and will be removed in a future version. Please use `Model.evaluate`, which supports generators.
    inceptionv3_scores = inceptionv3_model.evaluate_generator(test_generator, steps=test_samples // BATCH_SIZE)
inceptionv3 Test Loss: 0.0008152585942298174
inceptionv3 Test Accuracy: 1.0

In [12]: # Get the number of test samples
num_test_samples = test_generator.n

# Calculate the number of steps for prediction
num_prediction_steps = num_test_samples // test_generator.batch_size + 1

# Generate predictions for all test samples
inceptionv3_predictions = inceptionv3_model.predict_generator(test_generator, steps=num_prediction_steps, verbose=1)

# Convert predictions to class labels
inceptionv3_predicted_labels = np.argmax(inceptionv3_predictions, axis=1)

C:\Users\mohit\AppData\Local\Temp\ipykernel_34676\2714831293.py:8: UserWarning: `Model.predict_generator` is deprecated and will be removed in a future version. Please use `Model.predict`, which supports generators.
    inceptionv3_predictions = inceptionv3_model.predict_generator(test_generator, steps=num_prediction_steps, verbose=1)
9/9 [=====] - 6s 613ms/step

In [13]: # Get true class labels
true_labels = test_generator.classes

In [14]: # Calculate classification report
inceptionv3_report = classification_report(true_labels, inceptionv3_predicted_labels)

In [15]: print("inception Classification Report:")
print(inceptionv3_report)

inception Classification Report:
      precision    recall  f1-score   support

      0         1.00      1.00      1.00       140
      1         1.00      1.00      1.00       140

   accuracy          1.00
  macro avg          1.00
 weighted avg          1.00

In [16]: inceptionv3_scores = inceptionv3_model.evaluate_generator(test_generator, steps=test_samples // BATCH_SIZE)
print("inceptionv3 Test Loss:", inceptionv3_scores[0])
print("inceptionv3 Test Accuracy:", inceptionv3_scores[1])

C:\Users\mohit\AppData\Local\Temp\ipykernel_34676\1195483549.py:1: UserWarning: `Model.evaluate_generator` is deprecated and will be removed in a future version. Please use `Model.evaluate`, which supports generators.
    inceptionv3_scores = inceptionv3_model.evaluate_generator(test_generator, steps=test_samples // BATCH_SIZE)
inceptionv3 Test Loss: 0.0008152585942298174
inceptionv3 Test Accuracy: 1.0

In [ ]:
```

```
In [1]: # Importing necessary libraries
import os
import random
import shutil
import numpy as np
import matplotlib.pyplot as plt
from keras.preprocessing.image import ImageDataGenerator
from keras.applications import VGG16, VGG19, ResNet50, InceptionV3, Xception
from keras.models import Model
from keras.layers import Dense, GlobalAveragePooling2D
from keras.optimizers import Adam
from keras.callbacks import ModelCheckpoint
from sklearn.metrics import classification_report, confusion_matrix

# Setting random seed for reproducibility
np.random.seed(42)
random.seed(42)

# Define constants
NUM_CLASSES = 2
IMG_SIZE = (224, 224)
BATCH_SIZE = 32
TRAIN_DIR = 'train'
VAL_DIR = 'val'
TEST_DIR = 'test'
TRAIN_SPLIT = 0.7
VAL_SPLIT = 0.1
TEST_SPLIT = 0.2
test_samples = 280

In [2]: # Define the base directory where the image data is located
BASE_DIR = r"C:\Users\mohit\Downloads\DIP Assignment 2\Data"

In [3]: # Set the paths for training, validation, and test data
TRAIN_DIR = os.path.join(BASE_DIR, 'train')
VAL_DIR = os.path.join(BASE_DIR, 'val')
TEST_DIR = os.path.join(BASE_DIR, 'test')

In [4]: # Load and preprocess the data
train_datagen = ImageDataGenerator(rescale=1./255, shear_range=0.2, zoom_range=0.2, horizontal_flip=True)
val_datagen = ImageDataGenerator(rescale=1./255)
test_datagen = ImageDataGenerator(rescale=1./255)

In [5]: train_generator = train_datagen.flow_from_directory(TRAIN_DIR, target_size=IMG_SIZE, batch_size=BATCH_SIZE, class_mode='categorical')

val_generator = val_datagen.flow_from_directory(VAL_DIR, target_size=IMG_SIZE, batch_size=BATCH_SIZE, class_mode='categorical')

test_generator = test_datagen.flow_from_directory(TEST_DIR, target_size=IMG_SIZE, batch_size=BATCH_SIZE, class_mode='categorical', shuffle=False)

Found 980 images belonging to 2 classes.
Found 140 images belonging to 2 classes.
Found 280 images belonging to 2 classes.

In [6]: # Define function for creating transfer learning models
def create_transfer_model(base_model, num_classes):
    x = base_model.output
    x = GlobalAveragePooling2D()(x)
    x = Dense(1024, activation='relu')(x)
    predictions = Dense(num_classes, activation='softmax')(x)
    model = Model(inputs=base_model.input, outputs=predictions)
    return model

# Define ResNet50 model
resnet50_base = ResNet50(include_top=False, weights='imagenet', input_shape=(IMG_SIZE[0], IMG_SIZE[1], 3))
resnet50_model = create_transfer_model(resnet50_base, NUM_CLASSES)

In [7]: # Compile the models
optimizer = Adam(learning_rate=0.0001)
resnet50_model.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=['accuracy'])

In [8]: # Define checkpoint callback to save the best model during training
checkpoint = ModelCheckpoint('best_model.h5', monitor='val_loss', save_best_only=True)

In [9]: # Train the models
resnet50_history = resnet50_model.fit_generator(train_generator, steps_per_epoch=train_generator.n // train_generator.batch_size, epochs=10, validation_data=val_generator, validation_steps=val_generator.n // val_generator.batch_size, callbacks=[checkpoint])

C:\Users\mohit\AppData\Local\Temp\ipykernel_33032\951620722.py:2: UserWarning: `Model.fit_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.
  resnet50_history = resnet50_model.fit_generator(train_generator, steps_per_epoch=train_generator.n // train_generator.batch_size, epochs=10, validation_data=val_generator, validation_steps=val_generator.n // val_generator.batch_size, callbacks=[checkpoint])
Epoch 1/10
30/30 [=====] - 358s 11s/step - loss: 0.0822 - accuracy: 0.9684 - val_loss: 3.9794 - val_accuracy: 0.5000
Epoch 2/10
30/30 [=====] - 324s 11s/step - loss: 0.0019 - accuracy: 1.0000 - val_loss: 1.4914 - val_accuracy: 0.4922
Epoch 3/10
30/30 [=====] - 326s 11s/step - loss: 0.0045 - accuracy: 1.0000 - val_loss: 7.1890 - val_accuracy: 0.4844
Epoch 4/10
30/30 [=====] - 325s 11s/step - loss: 0.0068 - accuracy: 0.9979 - val_loss: 34.1761 - val_accuracy: 0.5000
Epoch 5/10
30/30 [=====] - 358s 12s/step - loss: 0.0027 - accuracy: 0.9989 - val_loss: 14.5630 - val_accuracy: 0.5078
Epoch 6/10
30/30 [=====] - 366s 12s/step - loss: 6.6791e-04 - accuracy: 1.0000 - val_loss: 2.8473 - val_accuracy: 0.5156
Epoch 7/10
30/30 [=====] - 391s 13s/step - loss: 2.8256e-04 - accuracy: 1.0000 - val_loss: 1.7559 - val_accuracy: 0.5000
Epoch 8/10
30/30 [=====] - 395s 13s/step - loss: 0.0084 - accuracy: 0.9958 - val_loss: 13.6165 - val_accuracy: 0.4922
Epoch 9/10
30/30 [=====] - 402s 13s/step - loss: 0.0079 - accuracy: 0.9968 - val_loss: 14.6081 - val_accuracy: 0.4844
Epoch 10/10
30/30 [=====] - 357s 12s/step - loss: 0.0020 - accuracy: 1.0000 - val_loss: 20.5592 - val_accuracy: 0.5000

In [10]: # Evaluate the models on test data
resnet50_scores = resnet50_model.evaluate_generator(test_generator, steps=test_samples // BATCH_SIZE)
print("ResNet50 Test Loss:", resnet50_scores[0])
print("ResNet50 Test Accuracy:", resnet50_scores[1])

C:\Users\mohit\AppData\Local\Temp\ipykernel_33032\556035509.py:2: UserWarning: `Model.evaluate_generator` is deprecated and will be removed in a future version. Please use `Model.evaluate`, which supports generators.
  resnet50_scores = resnet50_model.evaluate_generator(test_generator, steps=test_samples // BATCH_SIZE)
ResNet50 Test Loss: 18.02694320678711
ResNet50 Test Accuracy: 0.546875

In [11]: # Get the number of test samples
num_test_samples = test_generator.n

# Calculate the number of steps for prediction
num_prediction_steps = num_test_samples // test_generator.batch_size + 1

# Generate predictions for all test samples
resnet50_predictions = resnet50_model.predict_generator(test_generator, steps=num_prediction_steps, verbose=1)

# Convert predictions to class labels
resnet50_predicted_labels = np.argmax(resnet50_predictions, axis=1)

C:\Users\mohit\AppData\Local\Temp\ipykernel_33032\2677680890.py:8: UserWarning: `Model.predict_generator` is deprecated and will be removed in a future version. Please use `Model.predict`, which supports generators.
  resnet50_predictions = resnet50_model.predict_generator(test_generator, steps=num_prediction_steps, verbose=1)
9/9 [=====] - 25s 3s/step

In [12]: # Get true class labels
true_labels = test_generator.classes

In [13]: # Calculate classification report
from sklearn.metrics import classification_report

# Get the ground truth labels
ground_truth_labels = test_generator.classes

# Get the predicted labels
resnet50_predicted_labels = np.argmax(resnet50_predictions, axis=1)

# Calculate classification report
classification_report_resnet50 = classification_report(ground_truth_labels, resnet50_predicted_labels, zero_division=1)

In [15]: print("ResNet50 Classification Report:")
print(classification_report_resnet50)

ResNet50 Classification Report:
              precision    recall  f1-score   support

         0            0.50         1.00         0.67         140
         1            1.00         0.00         0.00         140

 accuracy            0.75
 macro avg           0.75         0.50         0.33         280
weighted avg           0.75         0.50         0.33         280
```